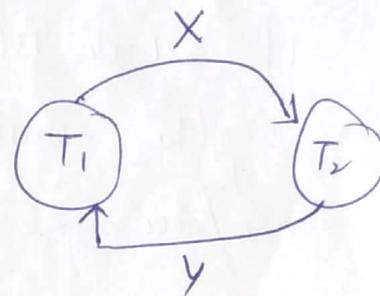
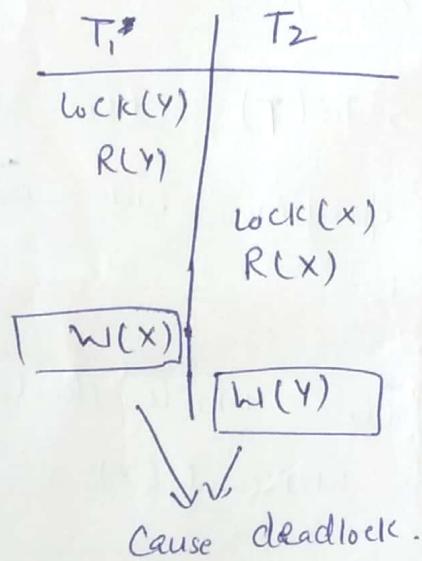


# # Dealing with Deadlocks

Deadlock: It occurs when each transaction  $T$  in a set of two or more transactions is waiting for some item that is locked by some other transaction  $T'$  in the set.

Eg:



## # Deadlock Prevention:

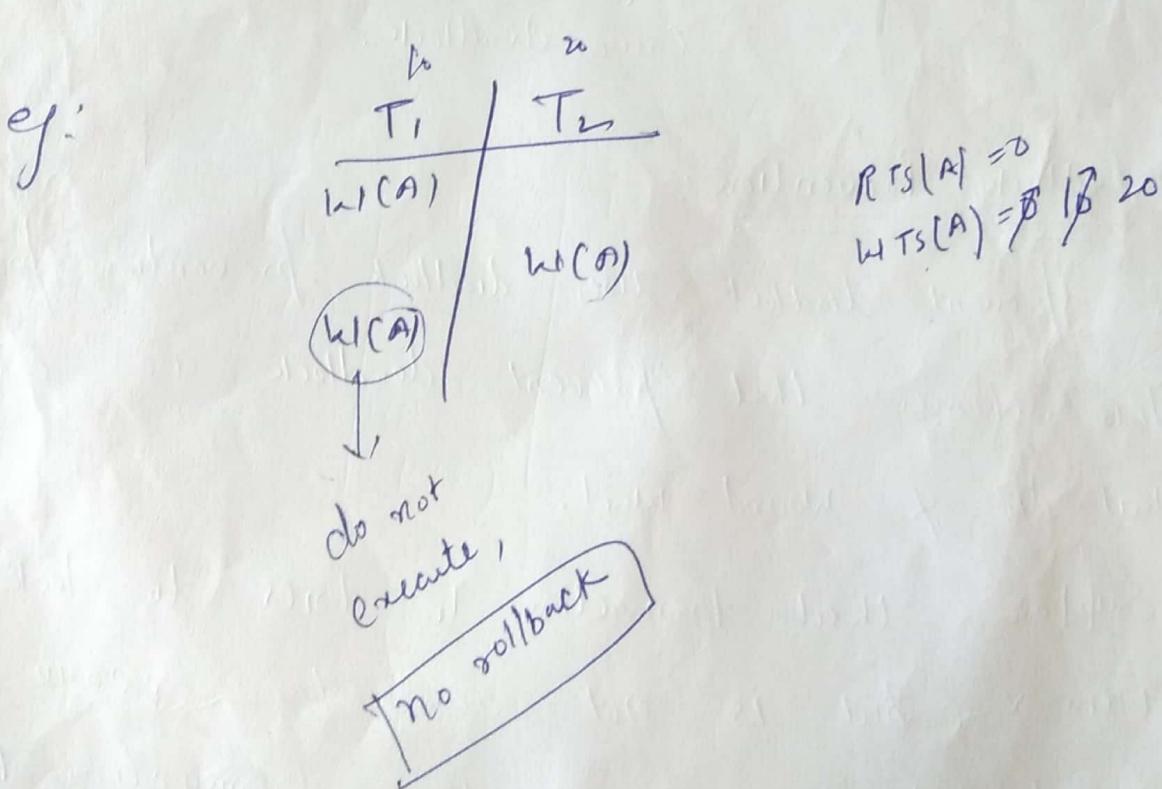
- To prevent deadlock, use deadlock prevention protocol.
- Two schemes that prevent deadlock are called wait-die & wound-wait.
- Suppose that transaction  $T_i$  tries to lock an item  $X$  but is not able to because  $X$  is locked by some other transaction  $T_j$  with a conflicting lock. The rules followed by these schemes are as follows:-

(21)

## # Thomas's Write Rule:

Modification of basic TO by modifying the check for the write-item( $x$ ) operation as follows:-

- (1) if  $\text{read-TS}(x) > \text{TS}(T)$ , then abort  $\cancel{x}$   
rollback T & reject the operation.
- (2) if  $\text{write-TS}(x) > \text{TS}(\cancel{T})$ , then do not execute the write operation but continue processing.
- (3) else execute the write-item( $x$ ) & set  $\text{write-TS}(x) = \text{TS}(T)$



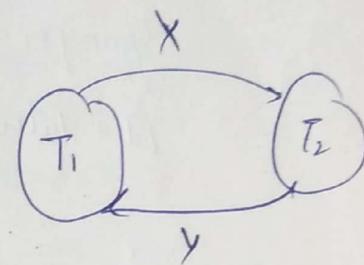
## # Deadlock

(22)

→ It occurs when each transaction  $T_i$  in a set of two or more transactions is waiting for some item that is locked by some other transaction  $T_j$  in the set.

e.g: →

$T_1$	$T_2$
dead-lock(y);	
dead-item(y);	
	read-lock(x);
	dead-item(x);
write-lock(x);	
	write-lock(y);



## # Deadlock Prevention Protocols

- One way to prevent deadlock, use deadlock prevention protocol.
- Two schemes that prevent deadlock are called wait-die & bounded-wait.
- Suppose that the transaction  $T_i$  tries to lock an item  $X$  but is not able to because  $X$  is locked by some other transaction  $T_j$  with a conflicting lock. The rules followed by these schemes are as follows:

1) kbit-die: It is a non-preemptive technique.

ii) If  $TS(T_i) < TS(T_j)$  that means  $T_i$  which is requesting a conflicting lock, is older than  $T_j$ , then  $T_i$  is allowed to wait.

ii) If  $TS(T_i) > TS(T_j)$  that means  $T_i$  is younger than  $T_j \rightarrow T_i$  dies and restart it later with the same timestamp.

(2)

	10	10	15
T <sub>1</sub>		T <sub>2</sub>	T <sub>3</sub>
lock(x)		lock(x) RE(x)	
			lock(x)

If  $T_1 > T_3$  both are requesting for  $x$

held by  $T_2$ ,

then  $T_1$  has to wait

&  $T_3 \rightarrow$  dies and rollback after some time.

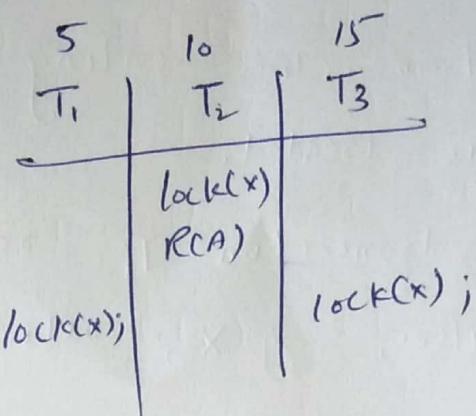
(2)

kbound-wait: It is a preemptive technique.

i) If  $TS(T_i) < TS(T_j)$ , then  $T_i$  forces  $T_j$  to be rolled back i.e.  $T_i$  wounds  $T_j$ .  $T_j$  is restarted later on with same timestamp.

2) If  $TS(T_i) > TS(T_j)$  then  $T_i$  is forced  
to wait until the resource is available. (Q3)

e.g:



$T_1 \& T_3$  are requesting for  $X$  hold by  $T_2$ ,  
then  $T_1$  waits  $T_2 \& T_2$  has to rolled  
back.

$\& T_3$  wait until database is available.

- # Deadlock detection / Avoidance :-
- In this, check whether the deadlock actually exists or not.
  - Deadlock is detected ~~so that~~ to check whether the deadlock is avoidable or not.
  - A simple way to detect a state of deadlock is for the system to construct and maintain wait-for graph.

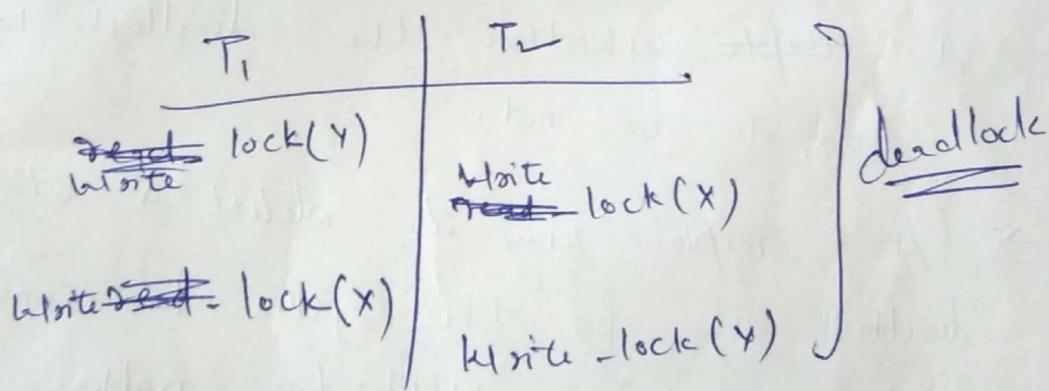
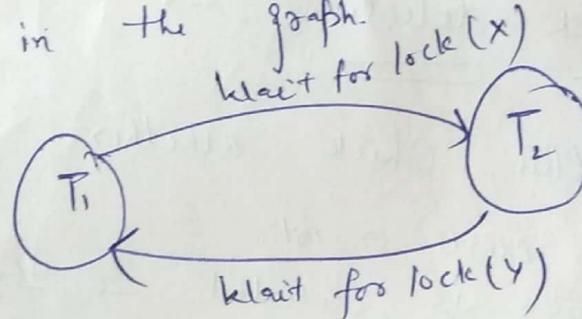
## # Wait-for-Graph : →

→ 1) For each transaction entering into the system, a node is created.

→ 2) When a transaction  $T_i$  requests for a lock on an item, say  $X$ , which is held by some other transaction  $T_j$ , a directed edge is created from  $T_i$  to  $T_j$ .

3) If  $T_j$  releases item  $X$ , the edge b/w them is dropped &  $T_i$  locks the data item.

4) The system keeps checking if there is any cycle in the graph.



The following approaches to avoid deadlock are:

(24)

- 1) Do not allow any request for an item, which is already locked by another transaction. This is not always feasible and may cause starvation, where a transaction indefinitely waits for a data item and can never acquire it.
- 2) Rollback one of the transactions. It is not always feasible to rollback the younger transaction, as it may be important than the older one. With the help of some relative algorithm, a transaction is chosen, which is to be aborted. This transaction is known as Victim. The process is known as Victim Selection.

Algorithm: for Victim Selection:

It should generally avoid selecting transactions that have been running for a long time and that have performed many updates, and it should try instead to select transactions that have not made any changes.

## # Introduction to Recovery Techniques →

- Recovery is required to protect the database from any data loss and inconsistencies.
- The recover manager ensures the atomicity of the ACID ~~durability~~ properties of a transaction.

## # Types of failures →

- 1) System crashes [eg: Power failure, Software Pbm, Virus etc]
- 2) System Errors [Pbm in software like syntax error eg: undesirable state] To avoid this, transaction has to rollback.
- 3) Disk failure
- 4) logical errors [eg: at a time of execution, correct data not found]  
[eg: do not identify identification]
- 5) Concurrency control enforcement [eg: when concurrent access on database, eg: lost update problem, Dirty read, Incorrect summary]
- 6) Natural & Physical failure [eg: power failure, Sparking, etc, earthquake, etc]

- System crash & logical errors are common type of failures & result in loss of volatile storage (such as main memory, cache) so they are also known as non-catastrophic failures.
- While, disk failure & failure due to nature result in loss of non-volatile storage & known as catastrophic failure.

Why Recovery is needed

CHITKARA  
UNIVERSITY



Format: CSE/Acad/01

Effective Date: 1-1-2017

Department of Computer Science & Engineering

Chitkara University Institute of Engineering & Technology

Types of failures: →

- ① Computer failure: H/w; S/w, n/w, crash.
- ② Transaction or system error
- ③ local errors or exception condition detected by transaction eg: no balance in account
- ④ Concurrency control enforcement : deadlock
- ⑤ Disk failure
- ⑥ Physical problems or catastrophic failures: power failure, fire, theft, etc.

## Recovery Techniques based on Deferred Update :- ①

→ The idea behind deferred update is to defer or postpone any actual updates to the database itself until the transaction completes its execution successfully & reaches its commit point.

The steps involved in the deferred update are as follows:-

1. When a transaction starts, write an entry start-transaction ( $T$ ) to the log;
2. When any operation is performed that will change values in the database, write a log entry write-item ( $T, X, \text{old-value}, \text{new-value}$ )
3. When a transaction is about to commit, write a log record of the form commit ( $T$ )
4. Commit the transaction, using the log to write the updates to database.
5. If transaction aborts, ignore the log records.  
Hence, there is no need to UNDO any operations.  
∴ This technique is also known as NO-UNDO/REDO

Algorithm.

- \* REDO is needed in case the system fails after the transaction commits but before all its changes are recorded in the database.

## #② Checkpoints :-

- Another type of entry in the log is called a ~~ch.~~ <sup>jj.</sup>
- A [checkpoint] record is written into the log periodically at that point when the system writes out to the database on disk all DBMS buffers that have been modified.
- All transactions that have their  $[commit, T]$  entries in the log before a [checkpoint] entry do not need to have their write operations redone in case of system crash, since all their updates will be recorded in the database on disk during checkpointing. [further from Navathe Pg = 675]

## # Recovery using Deferred update in a Single

### - User Environments :-

The algorithm RDUS uses a REDO procedure. →

### Algorithm RDUS :-

It uses two list of transactions:

- 1) Committed transactions since the last checkpoint
- 2) Active transactions

### Procedure REDO :-

Redoing a write-item operation WRITE-OP consist of examining its log entry and setting the value of item X in the database to new-value.

eg: [Start-transaction, T<sub>1</sub>]

[Write-item, T<sub>1</sub>, D, 26]

[commit, T<sub>1</sub>]

[start-transaction, T<sub>2</sub>]

[Write-item, T<sub>2</sub>, B, 10]

[Write-item, T<sub>2</sub>, D, 25]

→ system crash

∴ operations of T<sub>1</sub> & redone & T<sub>2</sub> log entries  
are ignored by recovery process.

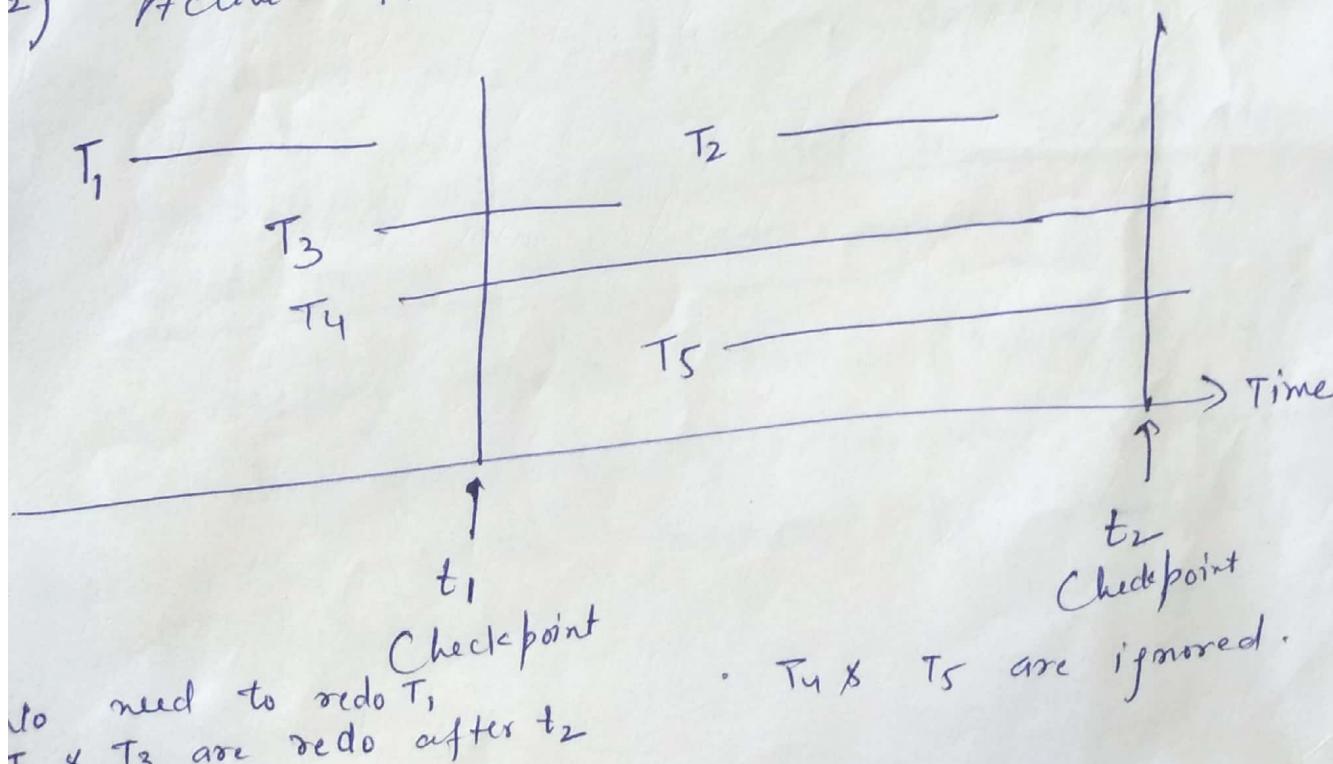
## # In Multiuser Environment :-

RDU-M uses REDO procedure:

Procedure RDU-M (with checkpoints) :-

Uses two list of transactions:

- 1) Committed transaction T since the last checkpoint  
(commit list)
- 2) Active transactions T' (active list)



is held by some

# REDO: Redo all the write operations of the transactions from the log, in the order in which they were written into the log. The transaction that are active and did not commit are effective, canceled and must be resubmitted.

e.g:

```
[start_transaction, T1]  
[write_item, T1, D1, 20]  
[commit, T1]  
[checkpoint]  
[start_transaction, T2]  
[write_item, T2, B1, 15]  
[commit, T2]  
[start_transaction, T3]  
[write_item, T3, B2, 12]
```

System crash

T<sub>2</sub> is ignored

T<sub>3</sub> is redone, T<sub>1</sub> is already saved.

## Recovery Techniques based on Immediate

(3)

Update :-

→ In these techniques, when a transaction issues an update command, the database can be updated immediately, without any need to wait for the transaction to reach its commit point.

Two main categories of immediate update algorithm are:-

(1) If the recovery technique ensures that all updates of a transaction are recorded in the database on disk before the transaction commits, there is never a need to REDO any operations of committed transactions. This is called as UNDO/ NO-REDO recovery algorithm.

(2) UNDO/ REDO recovery algorithm:

If transaction is allowed to commit before all its changes are written to database,

# UNDO/ REDO recovery based on immediate update in a Single-user Environment :-

The recovery algorithm RUV-s (recovery using immediate update in a Single-user environment) uses the REDO procedure as well as UNDO procedure:

Procedure RIU-S :→ (Pg-683)

- (1) Use two lists of transactions:  
Committed transactions & active transactions
- (2) Undo all write-item operations of active transactions.
- (3) Redo write-item operations of committed transactions.

Procedure UNDO :→ Undoing a write-item operation  
consists of examining its log entry.

Procedure REDO : Redoing a write-item operation.

# UNDO/REDO recovery based on immediate update with concurrent execution :→

Assume log includes checkpoints and concurrency control protocol produces strict schedules.

Procedure RIU-M :→

- 1) Use two lists of transactions: committed & active
- 2) Undo all write-item operations of active transactions
- 3) Redo write-item operations of committed "

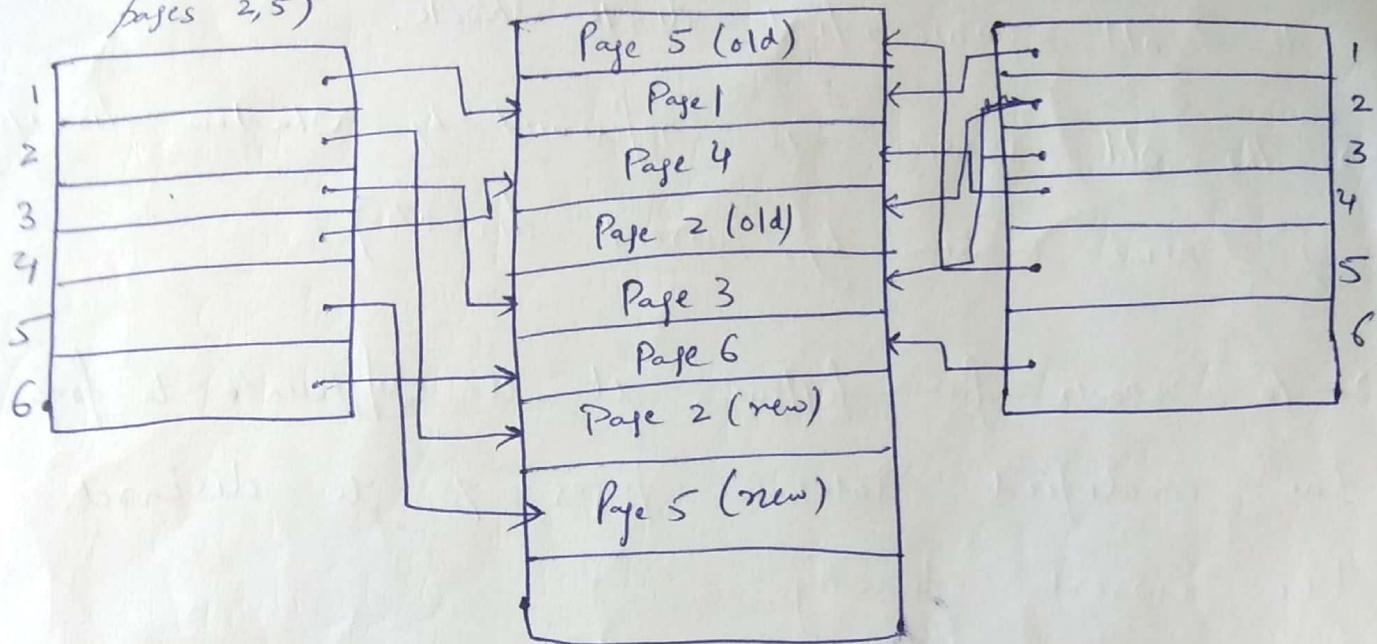
## Shadow Paging : $\rightarrow /(\text{NO-UNDO/NO-REDO})$

④

current directory  
(after updating  
pages 2,5)

Database  
disk  
blocks (pages)

Shadow directory  
(not updated)



- When a transaction begins executing, the current directory whose entries point to the most recent or current database pages on disk is copied into a shadow directory.
- The shadow directory is then saved on disk while the current directory is used by transaction.
- During transaction execution, shadow directory is never modified.
- When a write-item operation is performed, a new copy of the modified database page is created but the old copy of that page is not overwritten.

- The current directory entry is modified to ~~continues~~ whereas shadow to point to the ~~old~~ new block, whereas shadow directory is not modified & continues to point to old unmodified disk block.
- The old version is referenced by shadow directory & new version by current directory.
- To recover from failure, it is sufficient to free the modified database pages & to discard the current directory.
- The database is returned to state prior to transaction that was executing when the crash occurred & modified pages are discarded.
- Committing a transaction corresponds to discarding the previous shadow directory.

## # Introduction to Database Security :-

(b)

Database can be managed by 2 different modes of security controls:

- (1) Authentication
- (2) Authorization

### Authentication :-

- It is the process of confirming that a user logs in only in accordance with the rights to perform the activities he is authorized to perform.
- User Authentication can be performed at operating system level or database level itself.
- For authentication, it requires two different credentials, those are userid & password.

### Authorization :

- It is the process managed by database manager.
- The manager obtains information about current authenticated users.

Different ways of permissions available for authorization:

- 1) Primary permission : Grants the authorization ID directly.
- 2) Secondary " : Grants to groups if user is member.
- 3) Public " : Grants to all users.
- 4) Context-Sensitive " : Grants to trusted context role.

# Discretionary Access control based on Granting (Q) and Revoking Privileges :→

DAC: It is used to grant privileges to users including capability to access specific data files or fields in specific records.

There are two levels for assigning privileges to use the database system.

In case of DAC, a given user typically has different access rights.

(1) Account level :- DBA specifies the privileges that each account hold independently of other accounts.

It includes privileges like Create Table, Create View, Alter, Drop, Modify, Select, etc.

(2) Relation level :- At this level, DBA can control the privilege to access each individual relation or view in database.

→ The granting and revoking of privileges generally follow an authorization model for discretionary privileges known as Access Matrix Model,

where rows of a matrix  $M$  represent subjects (users) & columns represent objects (relations, records).

→ The following types of privileges can be granted on each individual relation  $R$  :-

1) Select privilege on  $R$

2) modify " " "  $R$ .

3) reference " " "  $R$ .

→ Revoking privileges : It is included for cancelling privileges.

# Privileges :- It represents the right of a particular user to access, create, manipulate & destroy various objects ~~inside~~ inside a database.

It can be granted in two ways:

- (1) Explicitly by giving object to user
- (2) By granting privileges to roles & then granting roles to users.

There are two distinct categories of privileges:

- (1) System Privilege: It allows user to perform some administrative tasks within a given DBMS. e.g.: Create Table, Session Create View  
Drop privileges.  
Alter any procedure, role.
- (2) Object Privilege: are those granted from a user to access to manipulate database objects.  
e.g.: Select, Insert, Update, Delete.

Ques:-

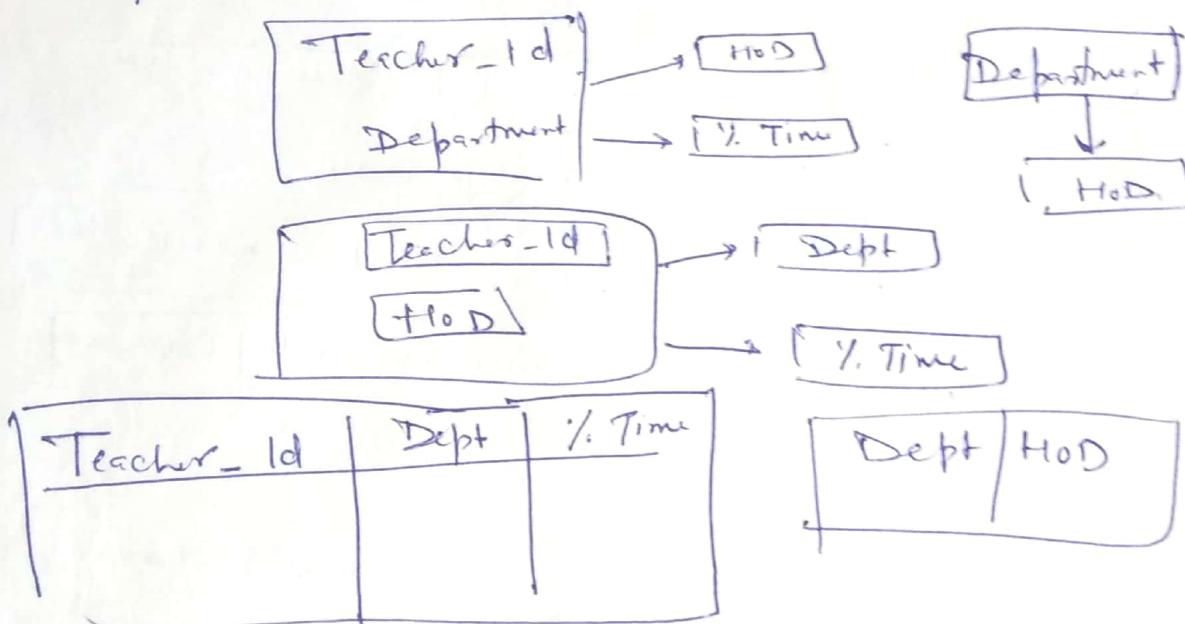
Teacher-Id	Department	HOD	Percent-Time
100	Computer	Vinod Kumbh	50
200	Mathematics	S. Sandhu	60
200	Physics	S. Pringal	40
300	History	Aruni Kumbh	30

Given Assume teacher can work in more than 1 dept,  
%age time he spent in each dept is given  
but each dept has only 1 HOD.

2 keys: → 1) (Teacher-Id, Department)

2) (Teacher-Id, HOD)

Dependencies:



Ques :-

Client\_interview      Relation :-

Client_no	Interview-date	Interview-time	Staff-no	Room-no
CR76	13/5/95.	10:30	SG5	9101
CR56	13/5/95	12:00	SG5	9101
CR74	13/5/95	12:00		
CR56	1/7/95	10:30	SG17 SG5	9102 9102

Dependencies :-

- 1) Client-no, Interview-date  $\rightarrow$  Interview-time, Staff-no, Room-no
  - 2) Staff-no, Interview-date  $\rightarrow$  Client-no Room-no
  - 3) Staff-no, Interview-date, Interview-time  $\rightarrow$  Client-no, Room-no
- $\downarrow$  BCNF

Client-no	Interview-date	Interview-time	Staffno
Staff-no	Interview-date	Room-no	

Ques

3NF :-

Order [ Order-no, Order-dt, Cust-name, Cust-addr,  
 Ship-addr, Phone1, Phone2, Cust-disc,  
 Product-code, Prod-desc, Qty, ~~Filled~~ Price ]

1) 1NF: Order [ Order-no, Order-dt, Cust-name, Cust-addr,  
 Ship-addr, Phone1 ]